

Project 2 MI3

Bharathi Thambidurai

2025-10-20

In MI2, we cleaned our data and stored the average minimum and maximum temperatures as time series objects in R. We also performed basic EDA to view how temperatures have changed over time, looking at specific months and looking at the overall trend. As we plotted our time series data, it was obvious there was a seasonal component we would have to keep in mind as we moved on to building a model.

```
# from MI2!!
mintemp <- read.csv("C:\\Users\\gbhar\\OneDrive - University of Virginia\\ds 4002\\mean_min_monthly_temp.csv")
maxtemp <- read.csv("C:\\Users\\gbhar\\OneDrive - University of Virginia\\ds 4002\\mean_max_monthly_temp.csv")

mintemp_dat <- mintemp[mintemp$Year != c(2024,2025), ]
maxtemp_dat <- maxtemp[maxtemp$Year != c(2024,2025), ]

# order of months
month_levels <- c("Jan", "Feb", "Mar", "Apr", "May", "Jun",
                  "Jul", "Aug", "Sep", "Oct", "Nov", "Dec")

# min temp reordered and stored as time series
mintemp_dat$Month <- factor(mintemp_dat$Month, levels = month_levels)
mintemp_sorted <- mintemp_dat[order(mintemp_dat$Year, mintemp_dat$Month), ]

mintemp_timeseries <- ts(mintemp_sorted$Mean_Min_Temp,frequency=12,start=c(2000,1))

# max temp reordered and stored as time series
maxtemp_dat$Month <- factor(maxtemp_dat$Month, levels = month_levels)
maxtemp_sorted <- maxtemp_dat[order(maxtemp_dat$Year, maxtemp_dat$Month), ]

maxtemp_timeseries <- ts(maxtemp_sorted$Value,frequency=12,start=c(2000,1))
```

We used <https://a-little-book-of-r-for-time-series.readthedocs.io/en/latest/src/timeseries.html> as a reference to guide us in our analysis.

Minimum Average Temperatures

We first built a model using exponential smoothing. Due to the obvious seasonality, we opted to use Holt-Winters exponential smoothing.

```
# fit predictive model
mintempmodel <- HoltWinters(mintemp_timeseries)
mintempmodel # estimated values of alpha, beta, gamma
```

```

## Holt-Winters exponential smoothing with trend and additive seasonal component.
##
## Call:
## HoltWinters(x = mintemp_timeseries)
##
## Smoothing parameters:
##   alpha: 0.1172071
##   beta  : 0.01432105
##   gamma: 0.1856539
##
## Coefficients:
##              [,1]
## a      47.940093133
## b       0.002574581
## s1  -18.105619060
## s2  -15.060639672
## s3   -9.520059313
## s4   -1.797080738
## s5    7.640912560
## s6   15.723107829
## s7   21.270803516
## s8   19.466157860
## s9   12.894557061
## s10   1.806070982
## s11 -10.017447454
## s12 -14.509994740


```

```

plot(mintempmodel) # plot model on top of existing data

```

Holt-Winters filtering



```
library(forecast)
```

```
## Warning: package 'forecast' was built under R version 4.4.3
```

```
## Registered S3 method overwritten by 'quantmod':
```

```
##   method          from
```


```
## as.zoo.data.frame zoo
```

```
# forecasting
```

```
mintempforecast <- forecast(mintempmodel, h=12) # 12 more months (2024)
```


```
plot(mintempforecast) # plot our forecast values
```

Forecasts from HoltWinters



```
# checking if there are any non-zero autocorrelations  
acf(na.omit(mintempforecast$residuals), lag.max=20) # correlogram
```

Series na.omit(mintempforecast\$residuals)



```
Box.test(na.omit(mintempforecast$residuals), lag=20, type = "Ljung-Box") # Ljung-Box test
```

```
##  
## Box-Ljung test  
##  
## data: na.omit(mintempforecast$residuals)  
## X-squared = 10.616, df = 20, p-value = 0.9556
```

Fitting our model and forecasting for 2024 shows that the model does pick up on seasonality and follow existing trends. To check our accuracy, we can see in the correlogram that autocorrelations for forecast errors do not exceed significance bounds for almost all lags. The p-value for our Ljung-Box test is 0.95, further proving there is little evidence of non-zero autocorrelations at lags 1-20.

We can also check if forecast errors have constant variance over time and are normally distributed with mean zero.


```
# from the guided tutorial we followed  
plotForecastErrors <- function(forecasterrors)  
{  
  forecasterrors <- na.omit(forecasterrors)  
  # make a histogram of the forecast errors:  
  mybinsize <- IQR(forecasterrors)/4  
  mysd <- sd(forecasterrors, na.rm=T)  
  mymin <- min(forecasterrors, na.rm=T) - mysd*5  
  mymax <- max(forecasterrors, na.rm=T) + mysd*5
```

```

# generate normally distributed data with mean 0 and standard deviation mysd
mynorm <- rnorm(10000, mean=0, sd=mysd)
mymin2 <- min(mynorm)
mymax2 <- max(mynorm)
if (mymin2 < mymin) { mymin <- mymin2 }
if (mymax2 > mymax) { mymax <- mymax2 }
# make a red histogram of the forecast errors, with the normally distributed data overlaid:
mybins <- seq(mymin, mymax, mybinsize)
hist(forecasterrors, col="red", freq=FALSE, breaks=mybins)
# freq=FALSE ensures the area under the histogram = 1
# generate normally distributed data with mean 0 and standard deviation mysd
myhist <- hist(mynorm, plot=FALSE, breaks=mybins)
# plot the normal curve as a blue line on top of the histogram of forecast errors:
points(myhist$mids, myhist$density, type="l", col="blue", lwd=2)
}

plot.ts(mintempforecast$residuals)

```




```

plotForecastErrors(mintempforecast$residuals) # normally distributed with mean 0

```

Histogram of forecasterrors



From the time plot, we can see that the forecast errors have constant variance over time. From the histogram, we can also assume that forecast errors are normally distributed with mean zero.

ARIMA MODEL

While Holt-Winters exponential smoothing provides an adequate predictive model, we can make a better model by taking correlations in the data into account. We wanted to try making an ARIMA (Autoregressive Integrated Moving Average) model as well.

Based on the tutorial and some trial-and-error in our model building, we realized our data needed to be differenced prior to determining what ARIMA parameters would best fit our model.

Once forcing a manual difference, we can use the `auto.arima` function to figure out the parameters of our ARIMA model.

```
mintemp_arima <- auto.arima(  
  mintemp_timeseries,  
  d = 1,           # force one regular difference  
  D = 1,           # seasonal difference (12-month)  
  max.p = 5, max.q = 5,  
  max.P = 2, max.Q = 2,  
  stepwise = FALSE, approx = FALSE,  
  trace = TRUE  
)  
  
##  
## ARIMA(0,1,0)(0,1,0)[12] : 1753.354  
## ARIMA(0,1,0)(0,1,1)[12] : Inf
```

```

## ARIMA(0,1,0)(0,1,2)[12] : 1589.917
## ARIMA(0,1,0)(1,1,0)[12] : 1665.576
## ARIMA(0,1,0)(1,1,1)[12] : 1590.103
## ARIMA(0,1,0)(1,1,2)[12] : Inf
## ARIMA(0,1,0)(2,1,0)[12] : 1643.012
## ARIMA(0,1,0)(2,1,1)[12] : 1591.984
## ARIMA(0,1,0)(2,1,2)[12] : Inf
## ARIMA(0,1,1)(0,1,0)[12] : Inf
## ARIMA(0,1,1)(0,1,1)[12] : Inf
## ARIMA(0,1,1)(0,1,2)[12] : Inf
## ARIMA(0,1,1)(1,1,0)[12] : 1513.352
## ARIMA(0,1,1)(1,1,1)[12] : Inf
## ARIMA(0,1,1)(1,1,2)[12] : Inf
## ARIMA(0,1,1)(2,1,0)[12] : 1490.029
## ARIMA(0,1,1)(2,1,1)[12] : Inf
## ARIMA(0,1,1)(2,1,2)[12] : Inf
## ARIMA(0,1,2)(0,1,0)[12] : Inf
## ARIMA(0,1,2)(0,1,1)[12] : Inf
## ARIMA(0,1,2)(0,1,2)[12] : Inf
## ARIMA(0,1,2)(1,1,0)[12] : 1515.409
## ARIMA(0,1,2)(1,1,1)[12] : Inf
## ARIMA(0,1,2)(1,1,2)[12] : Inf
## ARIMA(0,1,2)(2,1,0)[12] : 1491.978
## ARIMA(0,1,2)(2,1,1)[12] : Inf
## ARIMA(0,1,3)(0,1,0)[12] : Inf
## ARIMA(0,1,3)(0,1,1)[12] : Inf
## ARIMA(0,1,3)(0,1,2)[12] : Inf
## ARIMA(0,1,3)(1,1,0)[12] : 1517.363
## ARIMA(0,1,3)(1,1,1)[12] : Inf
## ARIMA(0,1,3)(2,1,0)[12] : 1494.027
## ARIMA(0,1,4)(0,1,0)[12] : Inf
## ARIMA(0,1,4)(0,1,1)[12] : Inf
## ARIMA(0,1,4)(1,1,0)[12] : 1519.449
## ARIMA(0,1,5)(0,1,0)[12] : Inf
## ARIMA(1,1,0)(0,1,0)[12] : 1661.144
## ARIMA(1,1,0)(0,1,1)[12] : Inf
## ARIMA(1,1,0)(0,1,2)[12] : Inf
## ARIMA(1,1,0)(1,1,0)[12] : 1587.862
## ARIMA(1,1,0)(1,1,1)[12] : Inf
## ARIMA(1,1,0)(1,1,2)[12] : Inf
## ARIMA(1,1,0)(2,1,0)[12] : 1567.045
## ARIMA(1,1,0)(2,1,1)[12] : Inf
## ARIMA(1,1,0)(2,1,2)[12] : Inf
## ARIMA(1,1,1)(0,1,0)[12] : Inf
## ARIMA(1,1,1)(0,1,1)[12] : Inf
## ARIMA(1,1,1)(0,1,2)[12] : Inf
## ARIMA(1,1,1)(1,1,0)[12] : Inf
## ARIMA(1,1,1)(1,1,1)[12] : Inf
## ARIMA(1,1,1)(1,1,2)[12] : Inf
## ARIMA(1,1,1)(2,1,0)[12] : Inf
## ARIMA(1,1,1)(2,1,1)[12] : Inf
## ARIMA(1,1,2)(0,1,0)[12] : Inf
## ARIMA(1,1,2)(0,1,1)[12] : Inf
## ARIMA(1,1,2)(0,1,2)[12] : Inf

```



```

## ARIMA(1,1,2)(1,1,0)[12] : Inf
## ARIMA(1,1,2)(1,1,1)[12] : Inf
## ARIMA(1,1,2)(2,1,0)[12] : 1494.057
## ARIMA(1,1,3)(0,1,0)[12] : Inf
## ARIMA(1,1,3)(0,1,1)[12] : Inf
## ARIMA(1,1,3)(1,1,0)[12] : Inf
## ARIMA(1,1,4)(0,1,0)[12] : Inf
## ARIMA(2,1,0)(0,1,0)[12] : 1630.335
## ARIMA(2,1,0)(0,1,1)[12] : Inf
## ARIMA(2,1,0)(0,1,2)[12] : Inf
## ARIMA(2,1,0)(1,1,0)[12] : 1562.511
## ARIMA(2,1,0)(1,1,1)[12] : Inf
## ARIMA(2,1,0)(1,1,2)[12] : Inf
## ARIMA(2,1,0)(2,1,0)[12] : 1532.239
## ARIMA(2,1,0)(2,1,1)[12] : Inf
## ARIMA(2,1,1)(0,1,0)[12] : Inf
## ARIMA(2,1,1)(0,1,1)[12] : Inf
## ARIMA(2,1,1)(0,1,2)[12] : Inf
## ARIMA(2,1,1)(1,1,0)[12] : Inf
## ARIMA(2,1,1)(1,1,1)[12] : Inf
## ARIMA(2,1,1)(2,1,0)[12] : 1494.019
## ARIMA(2,1,2)(0,1,0)[12] : Inf
## ARIMA(2,1,2)(0,1,1)[12] : Inf
## ARIMA(2,1,2)(1,1,0)[12] : 1519.455
## ARIMA(2,1,3)(0,1,0)[12] : Inf
## ARIMA(3,1,0)(0,1,0)[12] : 1620.793
## ARIMA(3,1,0)(0,1,1)[12] : Inf
## ARIMA(3,1,0)(0,1,2)[12] : Inf
## ARIMA(3,1,0)(1,1,0)[12] : 1550.529
## ARIMA(3,1,0)(1,1,1)[12] : Inf
## ARIMA(3,1,0)(2,1,0)[12] : 1523.49
## ARIMA(3,1,1)(0,1,0)[12] : Inf
## ARIMA(3,1,1)(0,1,1)[12] : Inf
## ARIMA(3,1,1)(1,1,0)[12] : 1519.443
## ARIMA(3,1,2)(0,1,0)[12] : Inf
## ARIMA(4,1,0)(0,1,0)[12] : 1614.582
## ARIMA(4,1,0)(0,1,1)[12] : Inf
## ARIMA(4,1,0)(1,1,0)[12] : 1545.551
## ARIMA(4,1,1)(0,1,0)[12] : Inf
## ARIMA(5,1,0)(0,1,0)[12] : 1604.366
##
##
## Best model: ARIMA(0,1,1)(2,1,0)[12]

```

```

mintemp_arima

```

```

## Series: mintemp_timeseries
## ARIMA(0,1,1)(2,1,0)[12]
##
## Coefficients:
##          ma1      sar1      sar2
##      -0.8623  -0.6417  -0.3145
## s.e.   0.0341   0.0598   0.0603

```

```
##
## sigma^2 = 12.62: log likelihood = -740.94
## AIC=1489.88 AICc=1490.03 BIC=1504.35
```


Based on the `auto.arima()` function, we can determine that an `ARIMA(0,1,1)x(2,1,0)[12]` best fits our data. We can use this model to forecast for 2024. We want to check that this model meets all assumptions.

```
mintempforecast2 <- forecast(mintemp_arima, h=12) # forecast using ARIMA model
mintempforecast2
```

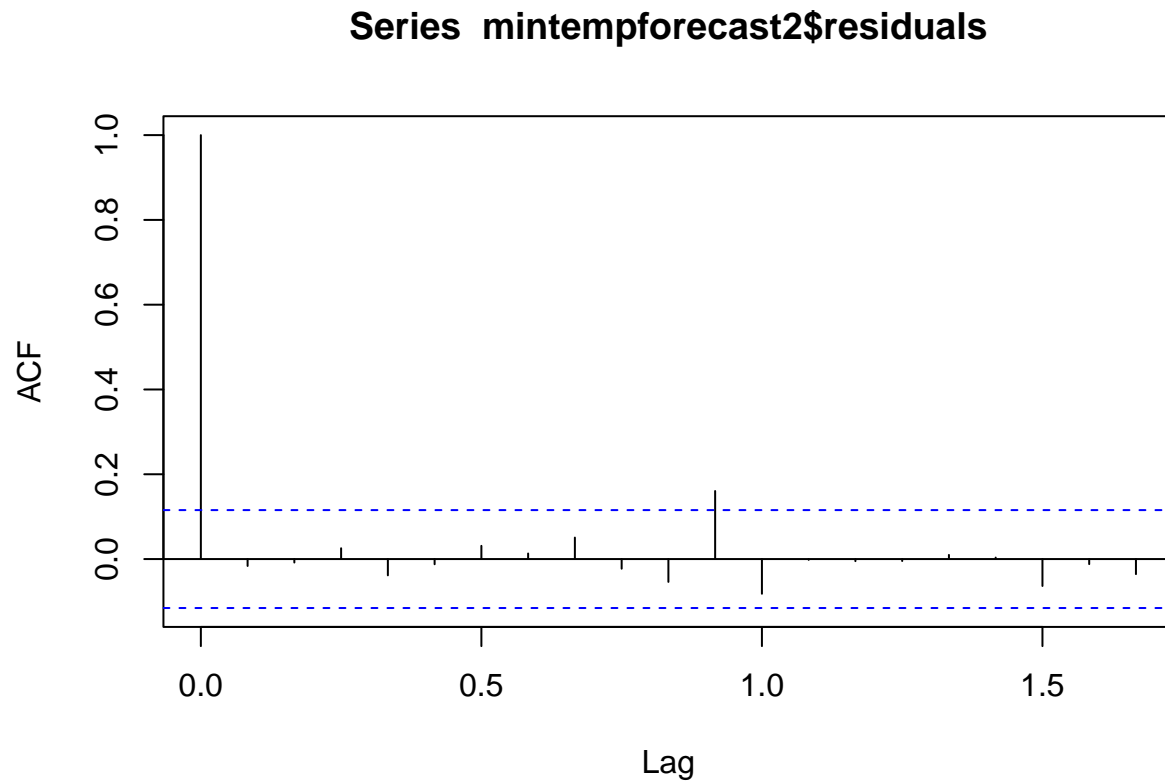
##		Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
##	Jan 2024	30.34521	25.79311	34.89731	23.38337	37.30705
##	Feb 2024	33.12620	28.53111	37.72128	26.09862	40.15377
##	Mar 2024	39.14963	34.51196	43.78729	32.05693	46.24232
##	Apr 2024	45.71244	41.03259	50.39230	38.55522	52.86967
##	May 2024	53.54706	48.82538	58.26873	46.32588	60.76824
##	Jun 2024	62.24026	57.47713	67.00338	54.95569	69.52483
##	Jul 2024	68.86638	64.06216	73.67060	61.51897	76.21380
##	Aug 2024	67.20413	62.35917	72.04909	59.79440	74.61386
##	Sep 2024	59.87884	54.99348	64.76421	52.40732	67.35036
##	Oct 2024	49.37681	44.45137	54.30225	41.84400	56.90962
##	Nov 2024	36.86023	31.89505	41.82542	29.26664	44.45383
##	Dec 2024	33.93048	28.92586	38.93510	26.27657	41.58438

```
plot(mintempforecast2) # add forecast values into time series plot
```

Forecasts from ARIMA(0,1,1)(2,1,0)[12]



```
acf(mintempforecast2$residuals, lag.max=20) # correlogram
```






```
Box.test(mintempforecast2$residuals, lag=20, type="Ljung-Box") # Ljung-Box test
```

```
##  
## Box-Ljung test  
##  
## data: mintempforecast2$residuals  
## X-squared = 14.434, df = 20, p-value = 0.8078
```


```
checkresiduals(mintemp_arima) # chcek that residuals are insignificant
```

Residuals from ARIMA(0,1,1)(2,1,0)[12]




```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(0,1,1)(2,1,0)[12]
## Q* = 23.707, df = 21, p-value = 0.3075
##
## Model df: 3.    Total lags used: 24
```

```
plot.ts(mintempforecast2$residuals)
```



```
plotForecastErrors(mintempforecast2$residuals) # normally distributed with mean 0
```

Histogram of forecasterrors



Based on our model and forecast values, we see that the correlogram shows that barely any of the sample autocorrelations for lags 1-20 exceed significance bounds, the p-value of the Ljung-Box test is 0.81, and the p-value of our residuals is 0.31. We can conclude there is little evidence for non-zero autocorrelations in the forecast errors.

From the time plot, we can see that the forecast errors have constant variance over time. From the histogram, we can also assume that forecast errors are normally distributed with mean zero.

Now, let's compare our forecast values from the ARIMA model to our actual values for 2024.

```
# extract 2024 from our data
mintemp_2024 <- mintemp[mintemp$Year==2024, ]
mintemp_2024
```

##	Month_Year	Year	Month	Mean_Min_Temp
## 25	Jan 2024	2024	Jan	30.9
## 51	Feb 2024	2024	Feb	34.5
## 77	Mar 2024	2024	Mar	42.8
## 103	Apr 2024	2024	Apr	51.4
## 129	May 2024	2024	May	58.9
## 155	Jun 2024	2024	Jun	65.4
## 181	Jul 2024	2024	Jul	68.1
## 207	Aug 2024	2024	Aug	65.0
## 233	Sep 2024	2024	Sep	58.5
## 259	Oct 2024	2024	Oct	46.2
## 285	Nov 2024	2024	Nov	40.8
## 311	Dec 2024	2024	Dec	30.1

```
mintempforecast2
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 2024	30.34521	25.79311	34.89731	23.38337	37.30705
## Feb 2024	33.12620	28.53111	37.72128	26.09862	40.15377
## Mar 2024	39.14963	34.51196	43.78729	32.05693	46.24232
## Apr 2024	45.71244	41.03259	50.39230	38.55522	52.86967
## May 2024	53.54706	48.82538	58.26873	46.32588	60.76824
## Jun 2024	62.24026	57.47713	67.00338	54.95569	69.52483
## Jul 2024	68.86638	64.06216	73.67060	61.51897	76.21380
## Aug 2024	67.20413	62.35917	72.04909	59.79440	74.61386
## Sep 2024	59.87884	54.99348	64.76421	52.40732	67.35036
## Oct 2024	49.37681	44.45137	54.30225	41.84400	56.90962
## Nov 2024	36.86023	31.89505	41.82542	29.26664	44.45383
## Dec 2024	33.93048	28.92586	38.93510	26.27657	41.58438

By comparing our forecast values of 2024 to our actual values of 2024, we can see that all the actual data falls in the 95% prediction interval for the forecast. Thus, we meet our goal of being at least 80% accurate with our model.

Maximum Average Temperatures


We again built a model using exponential smoothing. Due to the obvious seasonality, we opted to use Holt-Winters exponential smoothing.

```
# fit predictive model
maxtempmodel <- HoltWinters(maxtemp_timeseries)
maxtempmodel # estimated values of alpha, beta, gamma
```

```
## Holt-Winters exponential smoothing with trend and additive seasonal component.
##
## Call:
## HoltWinters(x = maxtemp_timeseries)
##
## Smoothing parameters:
##   alpha: 0.1484763
##   beta : 0.001741349
##   gamma: 0.2358454
##
## Coefficients:
##           [,1]
## a    72.78009694
## b     0.05615067
## s1 -23.05466568
## s2 -17.37428096
## s3 -10.47583962
## s4  -0.20418467
## s5   6.01131442
## s6  13.99878463
## s7  19.16614683
## s8  16.54661897
```


```
## s9 10.52109934
## s10 0.88235963
## s11 -10.66149562
## s12 -18.84281187
```

```
plot(maxtempmodel) # plot model on top of existing data
```




```
# forecasting
maxtempforecast <- forecast(maxtempmodel, h=12) # 12 more months (2024)
plot(maxtempforecast) # plot our forecast values
```


Forecasts from HoltWinters



```
# checking if there are any non-zero autocorrelations  
acf(na.omit(maxtempforecast$residuals), lag.max=20) # correlogram
```

Series na.omit(maxtempforecast\$residuals)




```
Box.test(na.omit(maxtempforecast$residuals), lag=20, type = "Ljung-Box") # Ljung-Box test
```

```
##  
## Box-Ljung test  
##  
## data: na.omit(maxtempforecast$residuals)  
## X-squared = 14.52, df = 20, p-value = 0.8032
```

Fitting our model and forecasting for 2024 shows that the model again does pick up on seasonality and follow existing trends. To check our accuracy, we can see in the correlogram that autocorrelations for forecast errors do not exceed significance bounds for all lags. The p-value for our Ljung-Box test is 0.80, further proving there is little evidence of non-zero autocorrelations at lags 1-20.


We can also check if forecast errors have constant variance over time and are normally distributed with mean zero.

```
plot.ts(maxtempforecast$residuals)
```



```
plotForecastErrors(maxtempforecast$residuals) # normally distributed with mean 0
```

Histogram of forecasterrors



From the time plot, we can see that the forecast errors have constant variance over time. From the histogram, we can also assume that forecast errors are normally distributed with mean zero. Despite the curve being a bit wonky, it generally follows the normal distribution curve.

ARIMA MODEL

While Holt-Winters exponential smoothing provides an adequate predictive model, we again wanted to make a better model by taking correlations in the data into account. We wanted to try making an ARIMA (Autoregressive Integrated Moving Average) model for maximum average temperatures as well.

Based on the tutorial and some trial-and-error in our model building, we again noticed our data needed to be differenced prior to determining what ARIMA parameters would best fit our model.

Once forcing a manual difference, we can use the `auto.arima` function to figure out the parameters of our ARIMA model.

```
maxtemp_arima <- auto.arima(  
  maxtemp_timeseries,  
  d = 1,           # force one regular difference  
  D = 1,           # seasonal difference (12-month)  
  max.p = 5, max.q = 5,  
  max.P = 2, max.Q = 2,  
  stepwise = FALSE, approx = FALSE,  
  trace = TRUE  
)
```

```
##  
## ARIMA(0,1,0)(0,1,0)[12] : 1811.063
```

```

## ARIMA(0,1,0)(0,1,1)[12] : Inf
## ARIMA(0,1,0)(0,1,2)[12] : Inf
## ARIMA(0,1,0)(1,1,0)[12] : 1761.423
## ARIMA(0,1,0)(1,1,1)[12] : Inf
## ARIMA(0,1,0)(1,1,2)[12] : Inf
## ARIMA(0,1,0)(2,1,0)[12] : 1724.567
## ARIMA(0,1,0)(2,1,1)[12] : 1657.14
## ARIMA(0,1,0)(2,1,2)[12] : 1652.784
## ARIMA(0,1,1)(0,1,0)[12] : 1683.043
## ARIMA(0,1,1)(0,1,1)[12] : Inf
## ARIMA(0,1,1)(0,1,2)[12] : Inf
## ARIMA(0,1,1)(1,1,0)[12] : 1624.995
## ARIMA(0,1,1)(1,1,1)[12] : Inf
## ARIMA(0,1,1)(1,1,2)[12] : Inf
## ARIMA(0,1,1)(2,1,0)[12] : 1592.023
## ARIMA(0,1,1)(2,1,1)[12] : Inf
## ARIMA(0,1,1)(2,1,2)[12] : 1525.06
## ARIMA(0,1,2)(0,1,0)[12] : Inf
## ARIMA(0,1,2)(0,1,1)[12] : Inf
## ARIMA(0,1,2)(0,1,2)[12] : Inf
## ARIMA(0,1,2)(1,1,0)[12] : 1626.913
## ARIMA(0,1,2)(1,1,1)[12] : Inf
## ARIMA(0,1,2)(1,1,2)[12] : Inf
## ARIMA(0,1,2)(2,1,0)[12] : 1594.074
## ARIMA(0,1,2)(2,1,1)[12] : Inf
## ARIMA(0,1,3)(0,1,0)[12] : Inf
## ARIMA(0,1,3)(0,1,1)[12] : Inf
## ARIMA(0,1,3)(0,1,2)[12] : Inf
## ARIMA(0,1,3)(1,1,0)[12] : 1627.682
## ARIMA(0,1,3)(1,1,1)[12] : Inf
## ARIMA(0,1,3)(2,1,0)[12] : 1594.704
## ARIMA(0,1,4)(0,1,0)[12] : Inf
## ARIMA(0,1,4)(0,1,1)[12] : Inf
## ARIMA(0,1,4)(1,1,0)[12] : 1629.772
## ARIMA(0,1,5)(0,1,0)[12] : Inf
## ARIMA(1,1,0)(0,1,0)[12] : 1748.582
## ARIMA(1,1,0)(0,1,1)[12] : Inf
## ARIMA(1,1,0)(0,1,2)[12] : Inf
## ARIMA(1,1,0)(1,1,0)[12] : 1701.23
## ARIMA(1,1,0)(1,1,1)[12] : Inf
## ARIMA(1,1,0)(1,1,2)[12] : Inf
## ARIMA(1,1,0)(2,1,0)[12] : 1662.546
## ARIMA(1,1,0)(2,1,1)[12] : Inf
## ARIMA(1,1,0)(2,1,2)[12] : 1593.679
## ARIMA(1,1,1)(0,1,0)[12] : Inf
## ARIMA(1,1,1)(0,1,1)[12] : Inf
## ARIMA(1,1,1)(0,1,2)[12] : Inf
## ARIMA(1,1,1)(1,1,0)[12] : 1626.938
## ARIMA(1,1,1)(1,1,1)[12] : Inf
## ARIMA(1,1,1)(1,1,2)[12] : Inf
## ARIMA(1,1,1)(2,1,0)[12] : Inf
## ARIMA(1,1,1)(2,1,1)[12] : Inf
## ARIMA(1,1,2)(0,1,0)[12] : Inf
## ARIMA(1,1,2)(0,1,1)[12] : Inf

```

```

## ARIMA(1,1,2)(0,1,2)[12] : Inf
## ARIMA(1,1,2)(1,1,0)[12] : Inf
## ARIMA(1,1,2)(1,1,1)[12] : Inf
## ARIMA(1,1,2)(2,1,0)[12] : Inf
## ARIMA(1,1,3)(0,1,0)[12] : Inf
## ARIMA(1,1,3)(0,1,1)[12] : Inf
## ARIMA(1,1,3)(1,1,0)[12] : Inf
## ARIMA(1,1,4)(0,1,0)[12] : Inf
## ARIMA(2,1,0)(0,1,0)[12] : 1716.751
## ARIMA(2,1,0)(0,1,1)[12] : Inf
## ARIMA(2,1,0)(0,1,2)[12] : Inf
## ARIMA(2,1,0)(1,1,0)[12] : 1664.955
## ARIMA(2,1,0)(1,1,1)[12] : Inf
## ARIMA(2,1,0)(1,1,2)[12] : Inf
## ARIMA(2,1,0)(2,1,0)[12] : 1621.051
## ARIMA(2,1,0)(2,1,1)[12] : Inf
## ARIMA(2,1,1)(0,1,0)[12] : Inf
## ARIMA(2,1,1)(0,1,1)[12] : Inf
## ARIMA(2,1,1)(0,1,2)[12] : Inf
## ARIMA(2,1,1)(1,1,0)[12] : 1627.423
## ARIMA(2,1,1)(1,1,1)[12] : Inf
## ARIMA(2,1,1)(2,1,0)[12] : 1594.292
## ARIMA(2,1,2)(0,1,0)[12] : Inf
## ARIMA(2,1,2)(0,1,1)[12] : Inf
## ARIMA(2,1,2)(1,1,0)[12] : 1629.51
## ARIMA(2,1,3)(0,1,0)[12] : Inf
## ARIMA(3,1,0)(0,1,0)[12] : 1699.997
## ARIMA(3,1,0)(0,1,1)[12] : Inf
## ARIMA(3,1,0)(0,1,2)[12] : Inf
## ARIMA(3,1,0)(1,1,0)[12] : 1643.218
## ARIMA(3,1,0)(1,1,1)[12] : Inf
## ARIMA(3,1,0)(2,1,0)[12] : 1605.336
## ARIMA(3,1,1)(0,1,0)[12] : Inf
## ARIMA(3,1,1)(0,1,1)[12] : Inf
## ARIMA(3,1,1)(1,1,0)[12] : 1629.5
## ARIMA(3,1,2)(0,1,0)[12] : Inf
## ARIMA(4,1,0)(0,1,0)[12] : 1697.362
## ARIMA(4,1,0)(0,1,1)[12] : Inf
## ARIMA(4,1,0)(1,1,0)[12] : 1640.132
## ARIMA(4,1,1)(0,1,0)[12] : Inf
## ARIMA(5,1,0)(0,1,0)[12] : 1697.758
##
##
## Best model: ARIMA(0,1,1)(2,1,2)[12]

```

```
maxtemp_arima
```

```

## Series: maxtemp_timeseries
## ARIMA(0,1,1)(2,1,2)[12]
##
## Coefficients:
##          ma1    sar1    sar2    sma1    sma2
##      -0.8468  0.375  -0.1598  -1.3755  0.4570

```

```
## s.e.    0.0430  0.215   0.0744   0.2065  0.2054
##
## sigma^2 = 13.28:  log likelihood = -756.37
## AIC=1524.75   AICc=1525.06   BIC=1546.45
```


Based on the `auto.arima()` function, we can determine that an `ARIMA(0,1,1)x(2,1,2)[12]` best fits our data. We can use this model to forecast for 2024. We want to check that this model meets all assumptions.

```
maxtempforecast2 <- forecast(maxtemp_arima, h=12) # forecast using ARIMA model
maxtempforecast2
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 2024	49.40025	44.73044	54.07006	42.25839	56.54211
## Feb 2024	53.90951	49.18524	58.63379	46.68436	61.13467
## Mar 2024	61.29836	56.52025	66.07648	53.99086	68.60586
## Apr 2024	72.74837	67.91701	77.57973	65.35944	80.13730
## May 2024	80.60167	75.71765	85.48569	73.13220	88.07114
## Jun 2024	87.89314	82.95702	92.82926	80.34399	95.44229
## Jul 2024	92.58003	87.59235	97.56771	84.95203	100.20803
## Aug 2024	89.47833	84.43962	94.51704	81.77229	97.18438
## Sep 2024	83.84218	78.75296	88.93141	76.05888	91.62549
## Oct 2024	74.10981	68.97056	79.24906	66.25001	81.96961
## Nov 2024	61.50452	56.31573	66.69331	53.56895	69.44008
## Dec 2024	54.19398	48.95612	59.43184	46.18337	62.20459


```
plot(maxtempforecast2) # add forecast values into time series plot
```

Forecasts from `ARIMA(0,1,1)(2,1,2)[12]`



```
acf(maxtempforecast2$residuals, lag.max=20) # correlogram
```

Series maxtempforecast2\$residuals






```
Box.test(maxtempforecast2$residuals, lag=20, type="Ljung-Box") # Ljung-Box test
```

```
##  
## Box-Ljung test  
##  
## data: maxtempforecast2$residuals  
## X-squared = 13.791, df = 20, p-value = 0.8409
```


```
checkresiduals(maxtemp_arima) # chcek that residuals are insignificant
```


Residuals from ARIMA(0,1,1)(2,1,2)[12]




```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(0,1,1)(2,1,2)[12]
## Q* = 16.568, df = 19, p-value = 0.6191
##
## Model df: 5.   Total lags used: 24
```

```
plot.ts(maxtempforecast$residuals)
```



```
plotForecastErrors(maxtempforecast$residuals) # normally distributed with mean 0
```

Histogram of forecasterrors



Based on our model and forecast values, we see that the correlogram shows that barely any of the sample autocorrelations for lags 1-20 exceed significance bounds, the p-value of the Ljung-Box test is 0.84, and the p-value of our residuals is 0.62. We can conclude there is little evidence for non-zero autocorrelations in the forecast errors.

From the time plot, we can see that the forecast errors have constant variance over time. From the histogram, we can also assume that forecast errors are normally distributed with mean zero.

Now, let's compare our forecast values from the ARIMA model to our actual values for 2024.

```
# extract 2024 from our data
maxtemp_2024 <- maxtemp[maxtemp$Year==2024, ]
maxtemp_2024
```

```
##      Month_Year Year Month Value
## 25      Jan_2024 2024   Jan  48.5
## 51      Feb_2024 2024   Feb  56.2
## 77      Mar_2024 2024   Mar  64.7
## 103     Apr_2024 2024   Apr  72.4
## 129     May_2024 2024   May  79.4
## 155     Jun_2024 2024   Jun  89.2
## 181     Jul_2024 2024   Jul  89.3
## 207     Aug_2024 2024   Aug  86.2
## 233     Sep_2024 2024   Sep  76.2
## 259     Oct_2024 2024   Oct  71.9
## 285     Nov_2024 2024   Nov  61.8
## 311     Dec_2024 2024   Dec  46.6
```

maxtempforecast2

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 2024	49.40025	44.73044	54.07006	42.25839	56.54211
## Feb 2024	53.90951	49.18524	58.63379	46.68436	61.13467
## Mar 2024	61.29836	56.52025	66.07648	53.99086	68.60586
## Apr 2024	72.74837	67.91701	77.57973	65.35944	80.13730
## May 2024	80.60167	75.71765	85.48569	73.13220	88.07114
## Jun 2024	87.89314	82.95702	92.82926	80.34399	95.44229
## Jul 2024	92.58003	87.59235	97.56771	84.95203	100.20803
## Aug 2024	89.47833	84.43962	94.51704	81.77229	97.18438
## Sep 2024	83.84218	78.75296	88.93141	76.05888	91.62549
## Oct 2024	74.10981	68.97056	79.24906	66.25001	81.96961
## Nov 2024	61.50452	56.31573	66.69331	53.56895	69.44008
## Dec 2024	54.19398	48.95612	59.43184	46.18337	62.20459

By comparing our forecast values of 2024 to our actual values of 2024, we again can see that all the actual data falls in the 95% prediction interval for the forecast. Thus, we meet our goal of being at least 80% accurate with our model.

Using ARIMA models with differencing and a seasonality component allowed us to build accurate models which generated prediction intervals that managed to capture our actual data points.